

Homework 5: Logistic Regression

Instructions: Submit a single Jupyter notebook (.ipynb) of your work to Collab by 11:59pm on the due date. All code should be written in Python. **Be sure to show all the work involved in deriving your answers! If you just give a final answer without explanation, you may not receive credit for that question.**

You may discuss the concepts with your classmates, but write up the answers entirely on your own. Do not look at another student's answers, do not use answers from the internet or other sources, and do not show your answers to anyone. **Cite any sources you used outside of the class material (webpages, etc.), and list any fellow students with whom you discussed the homework concepts.**

Note: If you have late penalty waivers remaining, you may use them to effectively extend the deadline with full credit.

1. Implement a function to fit a logistic regression using gradient ascent to maximize the likelihood of the parameter β . Your implementation should take an $n \times d$ data matrix X and corresponding binary array y of n labels. As usual, you should implement this yourself, not using a library that does gradient ascent or logistic regression for you.

Use your function to fit a logistic regression to the OASIS hippocampus data from HW 1, with $d = 2$ features (left/right hippocampus volume) for x , and the **Dementia** diagnosis as the y labels. Use the same training and testing data split as before. Do the following:

- (a) Scale all of your x data so that both features are in the range $[0, 1]$. Put these into a data matrix X adding a column of 1's for the intercept. Randomly initialize your β estimate.
 - (b) You will have to **tune** the gradient ascent step size, δ . Pick an initial δ . Run your algorithm for a few iterations, watching the value of the log-likelihood. If it ever goes down, your step size is too big! Decrease it. If it goes up too slowly, your step size is too small! Increase it. Eventually you should settle on a step size that works well, and the algorithm will converge to gradient close to zero.
 - (c) Your function should save the log-likelihood values each iteration. Plot the log-likelihood (vertical axis) as a function of the iteration number (horizontal axis). This plot should be increasing and flatten out if your algorithm converges correctly.
 - (d) Plot a 2D scatterplot of the x training data points, with two different colors for healthy and dementia subjects. Then, draw a line corresponding to your classifier prediction $p(y = 1 | x) = 0.5$. (This is your classifier's estimated separating line between the two classes that we discussed in class.)
 - (e) Use your estimated logistic regression model to predict the **Dementia** diagnosis from the testing data x features. What is your final accuracy? How does it compare to the accuracy you got in HW 1 with naïve Bayes?
2. In this part you will run your logistic regression model on a higher-dimensional example, the CIFAR-10 dataset. This is a collection of 32×32 color images of 10 classes (bird, cat, dog, truck, etc.), which is often used in machine learning and computer vision research. Download the CIFAR-10 image data (Python version) from this website:

<https://www.cs.toronto.edu/~kriz/cifar.html>

Your goal is to build a classifier that can recognize the difference between images of airplanes and frogs.

- (a) Read in the file `data_batch_1`, containing 10,000 training data images and their class labels. Flatten images to be row vectors in a data matrix. Note: images are originally $3 \times 32 \times 32 = 3,072$, where you have 3 color channels and 32 pixels in width and height. In the end, you should have a data matrix of size $10,000 \times 3,072$. The corresponding class labels (each a digit from 0-9) will be a vector of size 10,000. **You may use the example code that is provided on the course webpage to do this step.**
- (b) Run your logistic regression fit using just the data with airplanes (class label 0) and frogs (class label 6). Repeat the same process as before to find a step size that gets your gradient ascent to converge.
- (c) Make the same plot of your log-likelihood function vs. the iteration number.
- (d) Read in the file `test_batch`, containing 10,000 test data images and their class labels. Use your estimated logistic regression model to predict the the labels of the airplanes and frogs images in the test data. Again, report your accuracy.
- (e) From your results, randomly pick ten airplanes and ten frogs examples that were classified correctly. Display these as a grid of images. Do the same for ten airplanes and ten frogs examples that were classified incorrectly. Do the mistakes that your classifier made seem reasonable (in other words, do you think these cases were more difficult)?